



hochschule mannheim

# Understanding Eventual Consistency

MSI Presentation SS2014

Horst Schneider, Patrick Beedgen

Hochschule Mannheim

June 17th, 2014

# Introduction

*„...the storage system guarantees that if no new updates are made to the object, eventually all accesses will return the last updated valuee“*  
–W. Vogels (2009)

# Introduction

## Interpretations of Eventual Consistency

### Interpretation 1:

*"When you read data[...], the response might not reflect the results of a recently completed write operation. The response might include some stale data. Consistency across all copies of the data is usually reached within a second; so if you repeat your read request after a short time, the response returns the latest data."*

### Interpretation 2:

*"This sort of system we term "single writer eventual consistency". So what are its properties?  
(1) A client could read stale data. (2) The client could see out-of-order write operations. [...] So this is our weakest form of consistency - eventually consistent with out of order reads in the short term."*

# Introduction

## Interpretations of Eventual Consistency

### DynamoDB Documentation

*"When you read data[...], the response might not reflect the results of a recently completed write operation. The response might include **some stale data**. Consistency across all copies of the data is **usually reached within a second**; so if you repeat your read request after a short time, the response returns the latest data."*

### MongoDB Documentation

*"This sort of system we term "single writer eventual consistency". So what are its properties?*  
*(1)A client could read stale data.*  
*(2)The client could see out-of-order write operations.[...]*  
*So this is our weakest form of consistency - eventually consistent with **out of order reads** in the short term."*

## The Problem

- Disparate and low-level formalisms  
*consistency model is tied to system implementation*
- Weak guarantees  
*in realistic scenarios updates **never** stop*
- Conflict resolution policies  
*resolution of conflicts in multiple replicas*
- Combinations of different consistency levels  
*strong consistency may be needed at certain times*

⇒ Some sort of formalism is needed to define semantics of Eventual Consistency

# Agenda

## ① Replicated Data Types

# Anfang Hauptteil Horst

## Replicated Data Types

- A replicated database stores **objects**  $\text{Obj} = \{x, y, \dots\}$



## Replicated Data Types

- A replicated database stores **objects**  $\text{Obj} = \{x, y, \dots\}$
- Every object  $x \in \text{Obj}$  has
  - a **value**  $\in \text{Val}$

## Replicated Data Types

- A replicated database stores **objects**  $\text{Obj} = \{x, y, \dots\}$
- Every object  $x \in \text{Obj}$  has
  - a **value**  $\in \text{Val}$
  - a **type**  $\text{type}(x)$

## Replicated Data Types

- A replicated database stores **objects**  $\text{Obj} = \{x, y, \dots\}$
- Every object  $x \in \text{Obj}$  has
  - a **value**  $\in \text{Val}$
  - a **type**  $\text{type}(x)$
  - **operations**  $\text{Op}_{\text{type}(x)}$  that a client can perform on it

## Replicated Data Types

- A replicated database stores **objects**  $\text{Obj} = \{x, y, \dots\}$
- Every object  $x \in \text{Obj}$  has
  - a **value**  $\in \text{Val}$
  - a **type**  $\text{type}(x)$
  - **operations**  $\text{Op}_{\text{type}(x)}$  that a client can perform on it
- Two examples: Int Register **intreg**, Counter **ctr**

$$\begin{aligned}\text{Op}_{\text{ctr}} &= \{\text{rd}, \text{inc}\} \\ \text{Op}_{\text{intreg}} &= \{\text{rd}, \text{wr}(k) \mid k \in \mathbb{Z}\}\end{aligned}$$

# Replicated Data Types

## Sequential Data Type Specification

In a *strongly consistent system*, the semantics of a data type can be described by a function

$$S_{\tau} : \text{Op}_{\tau}^{+} \rightarrow \text{Val}$$

# Replicated Data Types

## Sequential Data Type Specification

In a *strongly consistent system*, the semantics of a data type can be described by a function

$$S_{\tau} : \text{Op}_{\tau}^{+} \rightarrow \text{Val}$$

Examples:

$$S_{\text{ctr}}(\sigma_{\text{rd}}) = (\text{number of inc operations in } \sigma);$$

# Replicated Data Types

## Sequential Data Type Specification

In a *strongly consistent system*, the semantics of a data type can be described by a function

$$S_{\tau} : \text{Op}_{\tau}^{+} \rightarrow \text{Val}$$

Examples:

$$\begin{aligned} S_{\text{ctr}}(\sigma_{\text{rd}}) &= (\text{number of inc operations in } \sigma); \\ S_{\text{intreg}}(\sigma_{\text{rd}}) &= k; \text{ if } \text{wr}(0)\sigma = \sigma_1 \text{wr}(k)\sigma_2 \text{ and} \\ &\quad \sigma_2 \text{ does not contain wr operations} \end{aligned}$$

# Replicated Data Types

## Sequential Data Type Specification

In a *strongly consistent system*, the semantics of a data type can be described by a function

$$S_{\tau} : \text{Op}_{\tau}^{+} \rightarrow \text{Val}$$

Examples:

$$S_{\text{ctr}}(\sigma \text{rd}) = (\text{number of inc operations in } \sigma);$$

$$S_{\text{intreg}}(\sigma \text{rd}) = k; \text{ if } \text{wr}(0)\sigma = \sigma_1 \text{wr}(k)\sigma_2 \text{ and}$$


$\sigma_2$  does not contain wr operations

$$S_{\text{intreg}}(\sigma \text{wr}(k)) = S_{\text{ctr}}(\sigma \text{inc}) = \perp;$$




# Replicated Data Types

## Conflict Resolution Strategies



# Replicated Data Types


## Conflict Resolution Strategies



- ① Make concurrent operations commutative

# Replicated Data Types


## Conflict Resolution Strategies



- ① Make concurrent operations commutative
- ② Order concurrent operations

# Replicated Data Types


## Conflict Resolution Strategies



- ① Make concurrent operations commutative
- ② Order concurrent operations
- ③ Flag conflicts (let the user decide)

# Replicated Data Types

## Conflict Resolution Strategies



- 1 Make concurrent operations commutative
- 2 Order concurrent operations
- 3 Flag conflicts (let the user decide)
- 4 Resolve conflicts semantically

# Replicated Data Types

## Replicated Data Type Specification

- $S_\tau$  is not strong enough to formalize these strategies

# Replicated Data Types

## Replicated Data Type Specification

- $S_\tau$  is not strong enough to formalize these strategies
- visibility and order of preceding operations have to be included

# Replicated Data Types

## Replicated Data Type Specification

- $S_\tau$  is not strong enough to formalize these strategies
- visibility and order of preceding operations have to be included
- $F_\tau$ : takes an **operation context** and returns a **value**

$$F_\tau(C) \in \text{Val}$$



# Replicated Data Types

## Replicated Data Type Specification

- $S_\tau$  is not strong enough to formalize these strategies
- visibility and order of preceding operations have to be included
- $F_\tau$ : takes an **operation context** and returns a **value**

$$F_\tau(C) \in \text{Val}$$

- operation context  $C$  adds **visibility** and **arbitration relations** to preceding operations:

$$C = (f, V, \text{ar}, \text{vis})$$

# Replicated Data Types

## Replicated Data Type Specification

- $S_\tau$  is not strong enough to formalize these strategies
- visibility and order of preceding operations have to be included
- $F_\tau$ : takes an **operation context** and returns a **value**

$$F_\tau(C) \in \text{Val}$$

- operation context  $C$  adds **visibility** and **arbitration relations** to preceding operations:

$$C = (f, V, \text{ar}, \text{vis})$$
$$u \xrightarrow{\text{vis}} v, \text{vis} \subseteq V \times V$$

# Replicated Data Types

## Replicated Data Type Specification

- $S_\tau$  is not strong enough to formalize these strategies
- visibility and order of preceding operations have to be included
- $F_\tau$ : takes an **operation context** and returns a **value**

$$F_\tau(C) \in \text{Val}$$

- operation context  $C$  adds **visibility** and **arbitration relations** to preceding operations:

$$\begin{aligned} C &= (f, V, \text{ar}, \text{vis}) \\ u &\xrightarrow{\text{vis}} v, \text{vis} \subseteq V \times V \\ u &\xrightarrow{\text{ar}} v, \text{ar} \subseteq V \times V \end{aligned}$$



# Ende Hauptteil Horst

# Anfang Hauptteil Patrick

# Axiomatic Specification Framework

## Levels of Eventual Consistency

- With replicated data types we can define multiple forms of eventual consistency
  - Basic eventual consistency
  - Ordering guarantees
  - on-demand consistency strengthening
- Every form contains multiple axioms

# Axiomatic Specification Framework

## Client Interaction Model

- Clients often wish to perform multiple operations within some context
- bla



# Axiomatic Specification Framework

## Basic Eventual Consistency Axioms

- Axioms a database has to fulfill to be eventual consistent
- SOWF, ARWF, VISWF, RVAL, EVENTUAL, THINAIR

# Axiomatic Specification Framework

## Session guarantees

- Axioms that ensure that databases stay consistent within a single session with a client
- RYW, MR, WYRV, WFRA, MWV, MWA

# Axiomatic Specification Framework

## Causal Consistency Axioms

- POCV, POCA, COCV, COCA

# Ende Hauptteil Patrick

# Conclusion

- Which problems does the techreport solve?
- What is not solved by it?
- What do **we** think about it?